

Chpater 3: Introduction to SQL

CONTENTS

1 OVERVIEW OF SQL

- 1) Developed as *Sequel* by IBM as part of their System R project in the early 1970s. Evolved into SQL (Structured Query Language).
- 2) Various standards were published for SQL such as SQL-89, SQL-92, SQL:1999, SQL:2003, SQL:2006, SQL:2008, SQL:2011, and SQL:2016.
- 3) SQL has several parts:
 - a) **Data-definition Language (DDL):** Allows specification of
 - i) Defining relation schemas
 - ii) Deleting relations
 - iii) Modifying relation schemas
 - iv) Integrity constraints
 - v) View definitions
 - vi) Authorization
 - b) **Data Manipulation Language (DML):** Provides commands for
 - i) Querying information from databases
 - ii) Insert, delete and modify tuples in a database
 - c) **Transaction control:** Specify begin and end points of transactions
 - d) **Embedded/dynamic SQL:** Define how SQL can be embedded into general-purpose languages.
- 4) **Note:** Comments in SQL are written in between `/* ... */` or are preceded by `--`.
- 3) **int:** An integer. Full form **integer**. Machine-dependent size.
- 4) **smallint:** Small integer. Machine-dependent size.
- 5) **numeric(p, d):** A fixed-point number with a user-specified precision. Number has p digits and d of the p digits are to the right of the decimal point.
- 6) **real, double precision:** Floating-point and double-precision floating-point numbers. Machine-dependent precision.
- 7) **float(n):** A floating point number with at least n digit precision.

Each type includes a special **null** value, indicating an absent or unknown value. To store multilingual data using the Unicode representation, use the **nvar-char** type.

2.2 Basic Schema Definition

SQL relation is defined by **create table** command. Syntax:

```
create table tableName
(
    attr_1 domain_1,
    attr_2 domain_2,
    ...,
    attr_n domain_n,
    constr_1,
    constr_2,
    ...,
    constr_m
);
```

NOTE: ALL statements in SQL MUST end with a **semicolon (;)**.

Integrity constraints supported by SQL:

2.1 Basic Types

- 1) **char(n):** A fixed length character string of user-specified length n . Full form **character**. Padded with whitespace at the end, if any.
- 2) **varchar(n):** A variable-length character string of maximum length n specified by the user. Full form **character varying**.
- 1) **primary key($A_{j_1}, A_{j_2}, \dots, A_{j_m}$):** Specifies the attributes in brackets that form the primary-key for the relation. All attributes should be non-null and unique. Optional but recommended constraint.
- 2) **foreign key($A_{k_1}, A_{k_2}, \dots, A_{k_n}$) references s :** Values of attributes in brackets for any tuple

in the relation must correspond to primary key attributes of some tuple in relation s .

- 3) **not null**: Specifies that the null value is not allowed for that attribute. Excludes the **null** value from the domain of the attribute.

Note: Tuples that violate integrity constraints are not inserted in the relation.

To remove a relation from a database

- 1) Use **drop table r** . This command deletes all information about table r .
- 2) In contrast, **delete from r** does not delete the table r , but deletes all tuples in r .
- 3) After a table is dropped, the **create table** command must be used to insert tuples into it.

To alter the schema of an existing relation, the **alter table** command is used.

- 1) To add an attribute: **alter table r add A D** , where:
 - a) r is the relation to be altered.
 - b) A is the attribute to be added and D is its domain.
- 2) To remove an attribute: **alter table r drop A** .
 - a) r is the relation to be altered.
 - b) A is the attribute to be deleted from the relation.

3 BASIC STRUCTURE OF SQL QUERIES

Consists of three clauses: **select**, **from**, **where**.

- 1) Input relations specified in **from** clause.
- 2) Operations on input relations defined in **select** and **where** clauses.
- 3) Basic syntax:

```
select attr_1, attr_2, ..., attr_n
from table_1, table_2, ..., table_m
where P;
```

3.1 Queries on Relations

- 1) To specify attributes in the **select** clause, one can prepend the attribute name with the relation name separated by a dot (.).
- 2) In practice, removing duplicates is time-consuming. The default behaviour in SQL is to retain duplicates.
- 3) To force elimination of duplicates, **distinct** keyword can be used in the **select** clause.

To explicitly specify that duplicates should be retained, **all** is used in the **select** clause.

- 4) The resulting select clause looks like **select (distinct—all)**.
- 5) The **select** clause can also contain arithmetic expressions on attributes or constants. These queries do NOT modify the underlying relation.
- 6) The **where** clause selects those tuples that satisfy the given predicate. One can join multiple predicates using **and**, **or** and **not**.
- 7) The general meaning/result of an SQL query is as follows:
 - a) Generate a cartesian product of the relations in the **from** clause.
 - b) Apply the predicates in the **where** clause.
 - c) Output the attributes of each tuple specified in the **select** clause.

4 ADDITIONAL BASIC OPERATIONS

4.1 The Rename Operation

Names cannot always be derived from a relation because:

- 1) Two relations in the **from** clause may have the same attribute names.
- 2) Arithmetic expressions do not even have a name.
- 3) One may want to change the attribute name even if it is derivable to something more convenient for future use.

To rename attributes, the **as** clause is used.

- 1) Syntax: r **as** s .
- 2) It can appear in both **select** and **from** clauses (to rename relations for convenience).
- 3) A useful case for **as** in the **from** clause is when a relation is compared to itself, in which case the relation is given two names.
- 4) The identifiers used to rename relations are called **correlation names** in the SQL standard. Other names are **table alias**, **correlation variable**, **tuple variable**.

4.2 String Operations

- 1) Strings in SQL are enclosed in single or double quotes (" or ").
- 2) If a string contains one quote, we can use the other quote to enclose it.
- 3) Functions on strings in SQL:

- a) *Concatenation* using `||`.
 - b) Convert to uppercase using `upper(s)`.
 - c) Convert to lowercase using `lower(s)`.
 - d) Remove spaces at the end of the string using `trim(s)`.
- 4) Pattern matching in SQL using the **like** or **not like** operator. Special characters used are:
 - a) Percent (%) to match any substring.
 - b) Underscore (_) to match any character.
 - 5) Special characters can be escaped for literal usage, such as `\%`, `\\`, and so on.

4.3 Attribute Specification in the Select Clause

- 1) An asterisk (*) can be used in the **select** clause to denote *all attributes*.
- 2) For all attributes in a certain table, use `tableName.*`. For all attributes in *all* tables, simply use `*`.

4.4 Ordering the Display of Tuples

- 1) **order by** clause in SQL allows the ordering of tuples in a relation.
- 2) Syntax (in **where** clause): `... order by A1, A2, ...`
- 3) Default behaviour is ascending order. To specify the sort order per attribute, we use **order by A₁ (asc—desc), A₂ (asc—desc), ...**

4.5 Where-Clause Predicates

- 1) For simplifying **where** clauses, SQL includes a **(not) between** clause.
- 2) Syntax (in **where** clause): `... (not) between v1 and v2 ...`
- 3) Note that both end values are *inclusive* for this clause.
- 4) Multiple equalities can be combined by using a *row constructor* () instead of using **and**.

5 SET OPERATIONS

SQL supports three set operations: **union**, **intersect** and **except**.

5.1 The Union Operation

- 1) Syntax: `r1 union r2`, where
 - a) **union** denotes the union operation.
 - b) `r1` and `r2` are compatible relations or subqueries yielding compatible relations.

- 2) The **union** operation eliminates duplicates unlike the **select** clause. To retain duplicates, use **union all**.

5.2 The Intersect Operation

- 1) Syntax: `r1 intersect r2`, where
 - a) **intersect** denotes the except operation.
 - b) `r1` and `r2` are compatible relations or subqueries yielding compatible relations.
- 2) The **intersect** operation eliminates duplicates unlike the **select** clause. To retain duplicates, use **intersect all**.

5.3 The Except Operation

- 1) Syntax: `r1 except r2`, where
 - a) **except** denotes the except operation.
 - b) `r1` and `r2` are compatible relations or subqueries yielding compatible relations.
- 2) The **except** operation eliminates duplicates unlike the **select** clause. To retain duplicates, use **except all**.

6 NULL VALUES

- 1) Presents special problems while evaluating arithmetic operations and predicates.
- 2) The result of an arithmetic operation involving **null** is *always null*.
- 3) The result of a comparison involving **null** is *always unknown*.
- 4) Behaviour of **unknown** (T stands for True, F for False, and U for unknown):

v_1	v_2	and	or	not
T	U	U	T	F
F	U	F	U	T
U	U	U	U	U

- 5) If a **where** clause evaluates to **false** or **unknown** for a tuple, that tuple is not included in the resulting relation.
- 6) To test for a **null** value, SQL includes the **is (not) null** predicate. Similarly for **is (not) unknown**.

7 AGGREGATE FUNCTIONS

Aggregate Functions: Functions that take a collection of values as input and return a single value. SQL has five standard built-in aggregate functions:

Aggregate Function	SQL Function
Average	avg
Minimum	min
Maximum	max
Total	sum
Count	count

7.1 Basic Aggregation

- 1) The result of applying a simple aggregate function is a single attribute containing a single aggregate value.
- 2) We can rename the attribute to something more convenient using the **as** clause.
- 3) Usage of **distinct**:
 - a) Eliminate duplicates in aggregate functions. Use **distinct** in the aggregate expression.
 - b) The use of **distinct** with **count(*)** is **not** allowed.
 - c) Using **distinct** with **max** and **min** is legal, but makes no difference.
 - d) To specify duplicate retention, we use **all**, but it is the default behaviour.

7.2 Aggregation with Grouping

- 1) To group sets of tuples for applying aggregates, we use **group by** clause. Tuples with the same attribute on all values in the **group by** clause are placed in one group.
- 2) When **group by** is used, attributes that appear in the **select** clause without being aggregated **MUST** appear in the **group by** clause.

7.3 The Having Clause

- 1) Applies to groups constructed by the **group by** clause.
- 2) Comes after **group by** clause in a query.
- 3) Like the **select** clause, any attribute that is present in the **having** clause without aggregation **MUST** appear in the **group by** clause.
- 4) For a query containing aggregation, following are the rules to evaluate its result:
 - a) Evaluate the **from** clause to get a relation.
 - b) Apply the predicate of the **where** clause.
 - c) Group these tuples using the **group by** clause.
 - d) Apply the **having** clause to each group.
 - e) Project selected attributes mentioned in the **select** clause.

7.4 Aggregation with Null and Boolean Values

- 1) All aggregate functions except **count(*)** ignore null values in their input collection.
- 2) The **count** of an empty collection is defined to be 0. Other aggregate operations return **null** on an empty collection.
- 3) For **boolean** data types, **some** and **every** can be used to compute the disjunction (**or**) conjunction (**and**) of the values.

8 NESTED SUBQUERIES

A **subquery** is a **select-from-where** expression nested inside another SQL query. Use cases of subqueries are illustrated below.

8.1 Set Membership

- 1) (**not**) **in** connective tests for set membership.
- 2) Can also use these operators in enumerated sets, such as (val1, val2, ...).
- 3) Use row constructors to test set membership in another relation (if supported by the database system).

8.2 Set Comparison

- 1) To check if an element returns **true** on comparison with at least one element of another relation, use (**op**) **some**, where **op** is a comparison operator.
- 2) Note that **=some** is *identical* to **in**, while **<>some** is *not identical* to **not in**.
- 3) To check if an element returns **true** on comparison with all elements of another relation, use (**op**) **all**, where **op** is the comparison operator.
- 4) Note that **<>some** is *identical* to **not in**, while **=all** is *not identical* to **in**.

8.3 Test for Empty Relations

SQL provides an **exists** construct that returns **true** if the argument relation (subquery) is nonempty. A **not exists** construct is also provided that mirrors this functionality.

8.4 Test for the Absence of Duplicate Tuples

SQL provides a **unique** construct for testing whether the argument subquery has no duplicate tuples. A **not unique** construct that mirrors this action is also provided.

8.5 Subqueries in the From Clause

- 1) A **select-from-where** query returns a relation, and hence can be nested in the **from** clause of another **select-from-where** query.
- 2) Such subqueries can also be given names using the **as** constructor (depends on SQL implementation).
- 3) Nested subqueries in the **from** clause cannot use correlation variables from other relations in the **from** clause.
- 4) Since SQL:2003 a subquery in the **from** clause that is prefixed by **lateral** can access other correlation variables.

8.6 The With Clause

- 1) The **with** clause provides a way to define a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- 2) Introduced in SQL:1999.

8.7 Scalar Subqueries

- 1) A subquery that returns only one tuple containing a single attribute is called a **scalar subquery**.
- 2) Can occur in **select**, **where** and **having** clauses.

8.8 Scalar Without a From Clause

- 1) Certain queries may have subqueries that use a **from** clause but the top-level query does not.
- 2) This may lead to an error in some systems. In such cases a *dummy* relation is provided to act as a *placeholder* in the **from** clause of the top level query, such as *dual* in Oracle.

9 MODIFICATION OF THE DATABASE

9.1 Deletion

- 1) SQL syntax:

```
delete from $r$
where $P$;
```

- 2) If the **where** clause is omitted, *all* tuples in the argument relation are deleted.
- 3) The **delete** command operates on *only* one relation.
- 4) ALL tuples in the relation are *tested* before deletion.

9.2 Insertion

- 1) To insert into a relation, a single tuple or a relation containing a set of tuples is specified.
- 2) SQL syntaxes:

```
insert into tableName(Schema) values (
    v_1, v_2, ..., v_n);
insert into tableName(Schema) <
    subquery>;
```

Note: The subquery must be chosen carefully. For example, queries like

```
insert into r select * from r;
```

can insert an infinite number of tuples into *r* if *r* does not have a primary key constraint.

- 3) One can also insert tuples containing **null** values, provided the attribute is declared in the relation schema as **not null**.

Updates

- The **update** statement is used to modify some values of the tuples. – SQL syntax :

```
<div class="syntax" style="width: 20em;">
update    r    set    (A1, A2, ..., An)    =
(v1, v2, ..., vn) where P;
</div>
where
```

- *r* is the relation to be updated. - *A* is the assignment statement. - *P* is the predicate based on which the tuples should be updated. - To have multiple updates occur simultaneously without one affecting the other, a **case** construct is provided.

The syntax for this is:

```
<div class="syntax" style="width: 11em;">
update    r    set    A = case emsp; emsp;
when    P1    then    v1 emsp; emsp; when    P2
then    v2 emsp; emsp; ... emsp; emsp;
when    Pn    then    vn emsp; emsp; else    v0
end;
</div>
where
```

- *r* is the relation. - *A* is the new attribute. - *P_i* are the predicates. - *v_i* are the values or expressions for *A* for each case. - We can also use scalar subqueries in **set** clause. The **case** construct permits us to handle nulls if they are present.